TWO APPS · ONE SYNTHESIS ENGINE · NOTHING ELSE SHARED

Bed
& Pedal

A tool that lets someone who loves sound but cannot compose lead a room. And a drone-first practice tool that keeps an improviser out of producer mode.

Status   Bed: engine built & verified, app unbuilt. Pedal: shipped, unverified on speakers.

Repo   adamelfersmusic-web/AdamVaultOS

Files   sound-bath-app/ · public/looper/

This doc   self-contained. Works offline. Prints to PDF.

ORIENTATION

Two separate products. They share a synthesis approach and nothing else, and should stay that way.

01 Bed — the sound bath app

02 Why it exists

03 Who it's for

04 The signature story

05 What exists today

06 How the sound works

07 The domain facts

08 How the app works

09 Marketing

10 Business

11 Not building

12 Roadmap

13 Pedal — the practice looper

14 What's built

15 Roadmap

16 The competitive landscape

17 Where everything lives

PART ONE

Bed

A generative sound engine and conducting tool for sound bath practitioners — the people who play crystal singing bowls and lead group sessions.

Why it exists

> A person who loves sound but can't compose becomes someone who can lead a room.
>
> Everything is downstream of this sentence. If a feature doesn't serve it, cut it.

Leading an immersive sound bath currently requires four unrelated skills:

Playing bowls and gongs	**Common.** Many people have this.
Charisma, holding a group	**Common.** Yoga-teacher types have this.
Sound design in a DAW	**Rare.**
Conducting an ensemble through an arc	**Rare.**

Almost nobody has all four, so the ceiling on who can lead is set by the two that have nothing to do with why anyone got into this. People with genuine love for the practice, real instruments and real presence are locked out because they can't open Ableton and can't count an ensemble through a form.

Meanwhile the audio they'd actually produce is a **thin, stable palette** — drones, pads, a felt-piano motif, ocean, rain, crickets, chimes. That's confirmed by observation: the strongest practitioner in the region uses a handful of elements, reused, plus live improvisers and a light show. It emphatically does not need a DAW. It needs the right five things, tuned to the room.

THE INSIGHT UNDERNEATH

The sound design was never the art. Leading the room is. This removes the part that was blocking people from the part they're actually good at. *Never* frame it as "your audience can't tell the difference" — it's true, and it reads as contempt for the craft to exactly the people whose endorsement you need.

Who it's for

The bowl player who wants to lead	**Primary.** Owns a crystal or Tibetan set. Has played in someone else's bath. Loves this deeply, cannot compose, will not learn a DAW. Routinely spends $200–$2000 on a single bowl. Currently capped at "one person with bowls."
The yoga or meditation teacher	**Secondary.** Great voice and presence, already leads rooms. Uses a Spotify playlist behind sessions and knows it's the weak part. Needs a bed that ducks out of the speech band and follows their pacing rather than the reverse.
The solo clinical practitioner	**Tertiary.** Hospice, palliative care, integrative medicine. Needs the stripped version — pick a bed, press play, play a bowl over it. Note hospitals draw a hard line between board-certified music therapists and sound practitioners; the door is usually hospice or integrative medicine, often volunteer-first.

The signature story

Thursday. Musicians you love are passing through town. There's no event planned. By Saturday there are sixty people on the floor, because you spent fifteen minutes building a structure and the rest of the evening being present.

> The promise isn't "better sound baths."
> It's *you can say yes on Thursday.*
>
> Which makes fifteen minutes a hard spec, not a figure of speech. If a first-timer can't reach a runnable 50-minute structure in fifteen minutes, the positioning collapses. Test it with real people early.

What exists today

Two working prototypes. Both single self-contained HTML files, no dependencies. Both are *working instruments, not designed apps* — the interface is throwaway, the behaviour is the deliverable.

THE FOUR-BAND ENGINE — PROTOTYPE/BED.HTML

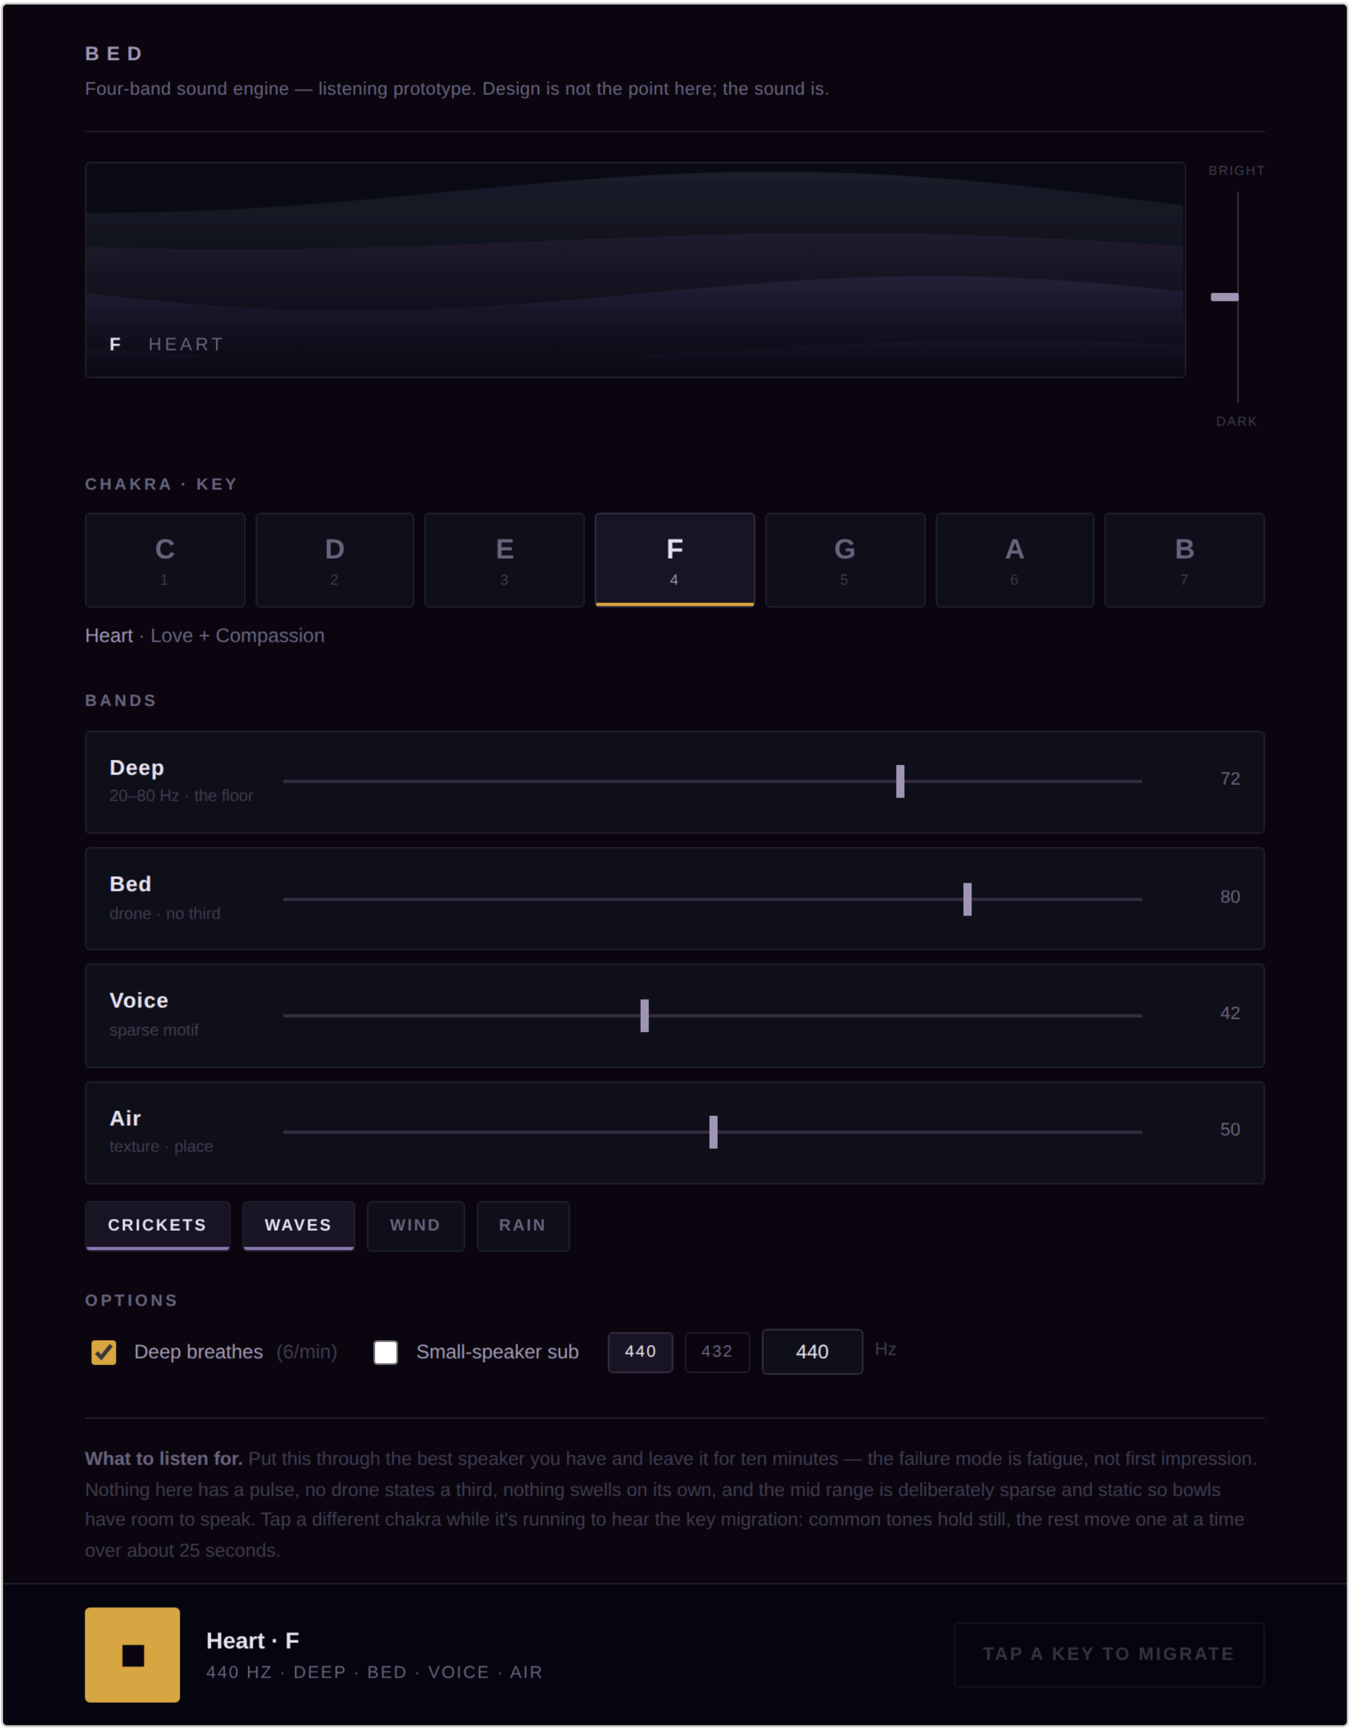

Four bands allocated by frequency so they cannot collide. Chakra selector derives the key. Tapping a different chakra while running migrates the whole room.

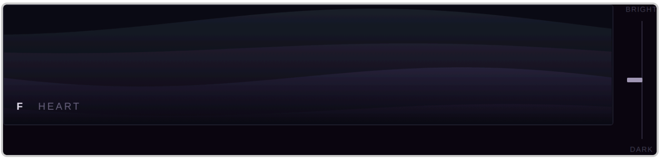

The four bands rendered as drifting fields.

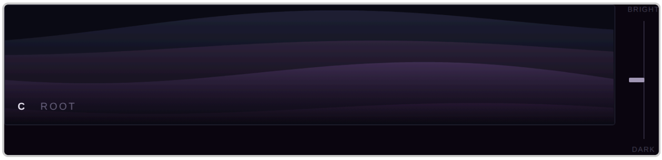

Mid-migration — the fields shift hue as the key travels.

ASSESSED BY A PRACTITIONER

“Already world class for a lot of people I know.” The sub confirmed as a genuine differentiator that nothing else in the space has. The frequency carving confirmed by singing a C bowl pitch over it and finding it sat correctly — the hole in the middle is the right size.

THE HANDOUT GENERATOR — `PROTOTYPE/SHEETS.HTML`

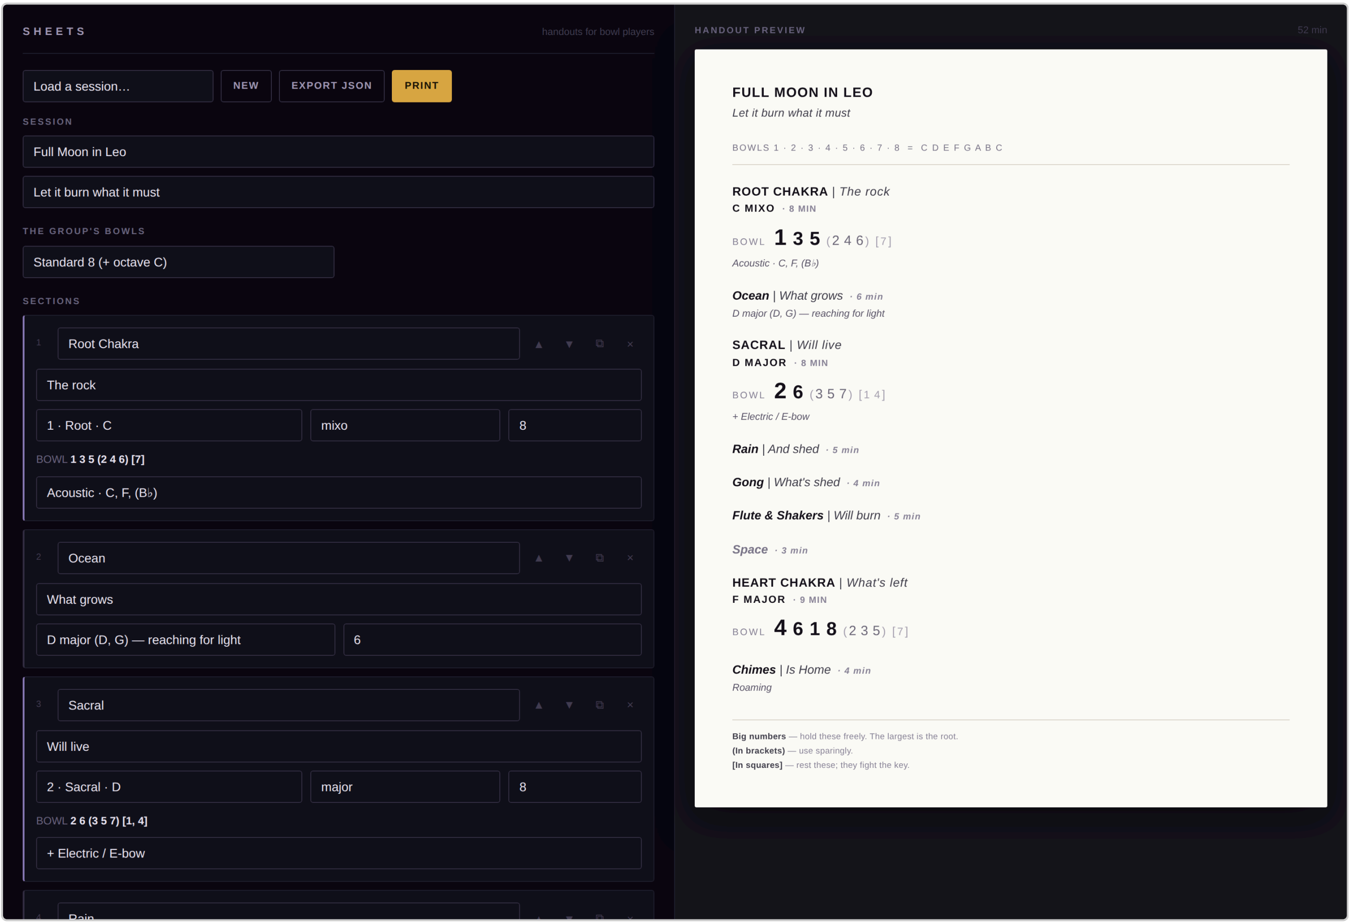

Build a session from chakra, mode and duration; the bowl notation computes live and prints black-on-white.

FULL MOON IN LEO

Let it burn what it must

BOWLS 1 · 2 · 3 · 4 · 5 · 6 · 7 · 8 = C D E F G A B C

ROOT CHAKRA | *The rock*

C MIXO · **8 MIN**

BOWL **1 3 5** (2 4 6) [7]

Acoustic · C, F, (B♭)

Ocean | *What grows* · ***6 min***

D major (D, G) — reaching for light

SACRAL | *Will live*

D MAJOR · **8 MIN**

BOWL **2 6** (3 5 7) [1 4]

+ Electric / E-bow

Rain | *And shed* · ***5 min***

Gong | *What's shed* · ***4 min***

Flute & Shakers | *Will burn* · ***5 min***

Space · ***3 min***

HEART CHAKRA | *What's left*

F MAJOR · **9 MIN**

BOWL **4 6 1 8** (2 3 5) [7]

Chimes | *Is Home* · ***4 min***

Roaming

Big numbers — hold these freely. The largest is the root.
(In brackets) — use sparingly.
[In squares] — rest these; they fight the key.

The printed handout. Tiers are carried by *size*, not vocabulary — readable at fifteen feet, in the dark, by someone who reads no music and possibly no English.

THE APP — `APP/INDEX.HTML`

The exploratory full build. One self-contained file, ~3,200 lines, no backend. Wizard to a runnable fifty-minute structure in about six taps.

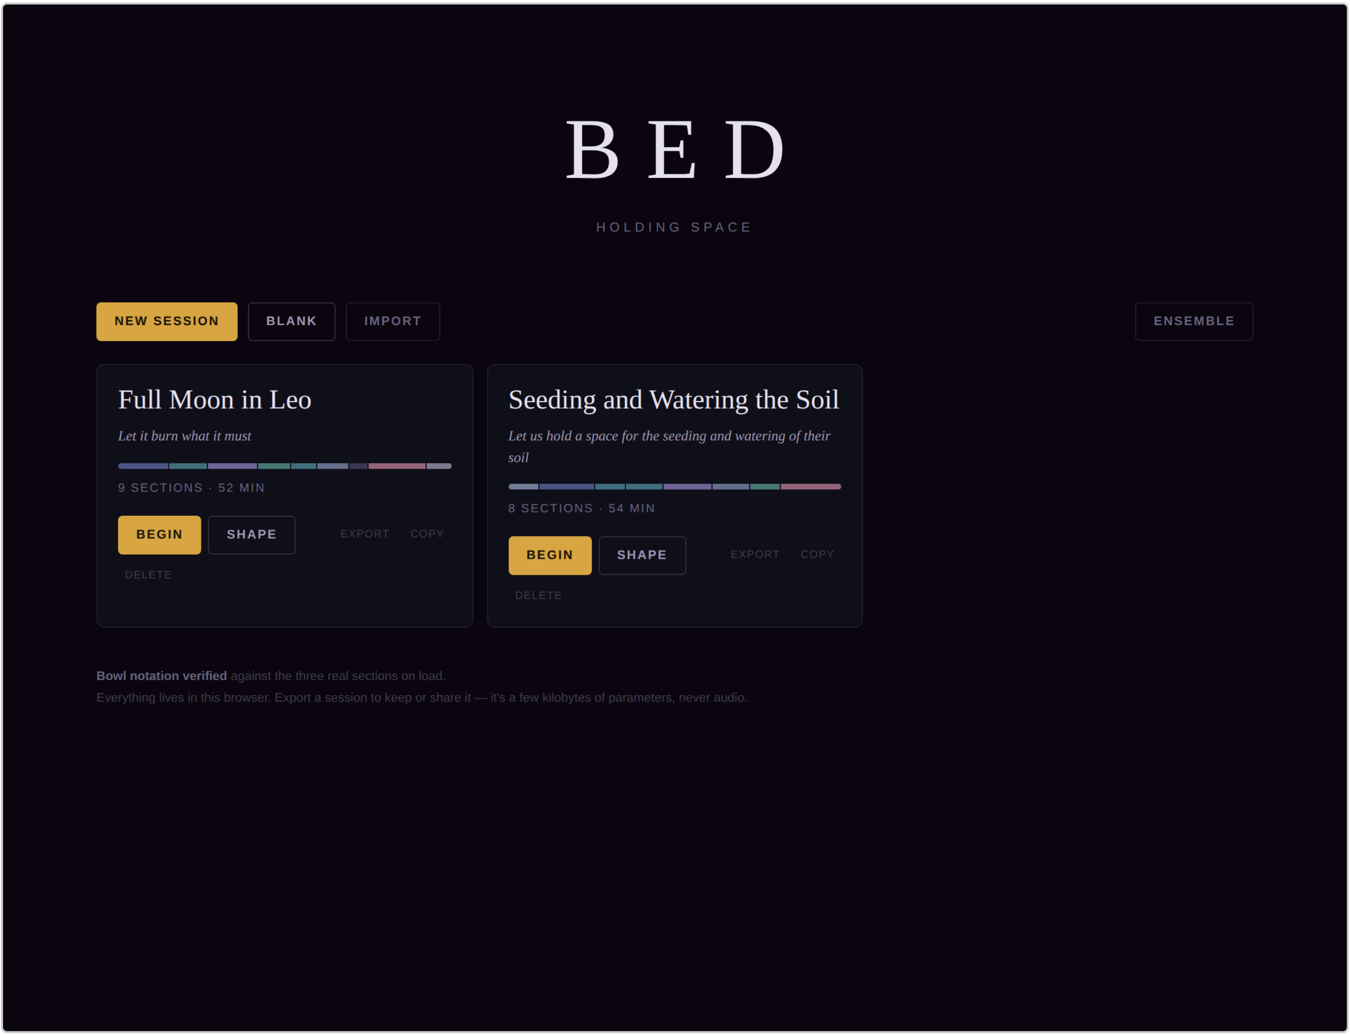

Library. Both real sessions seeded, each arc drawn as a colour strip. BEGIN and SHAPE as the two verbs — not play and edit, which would be software language.

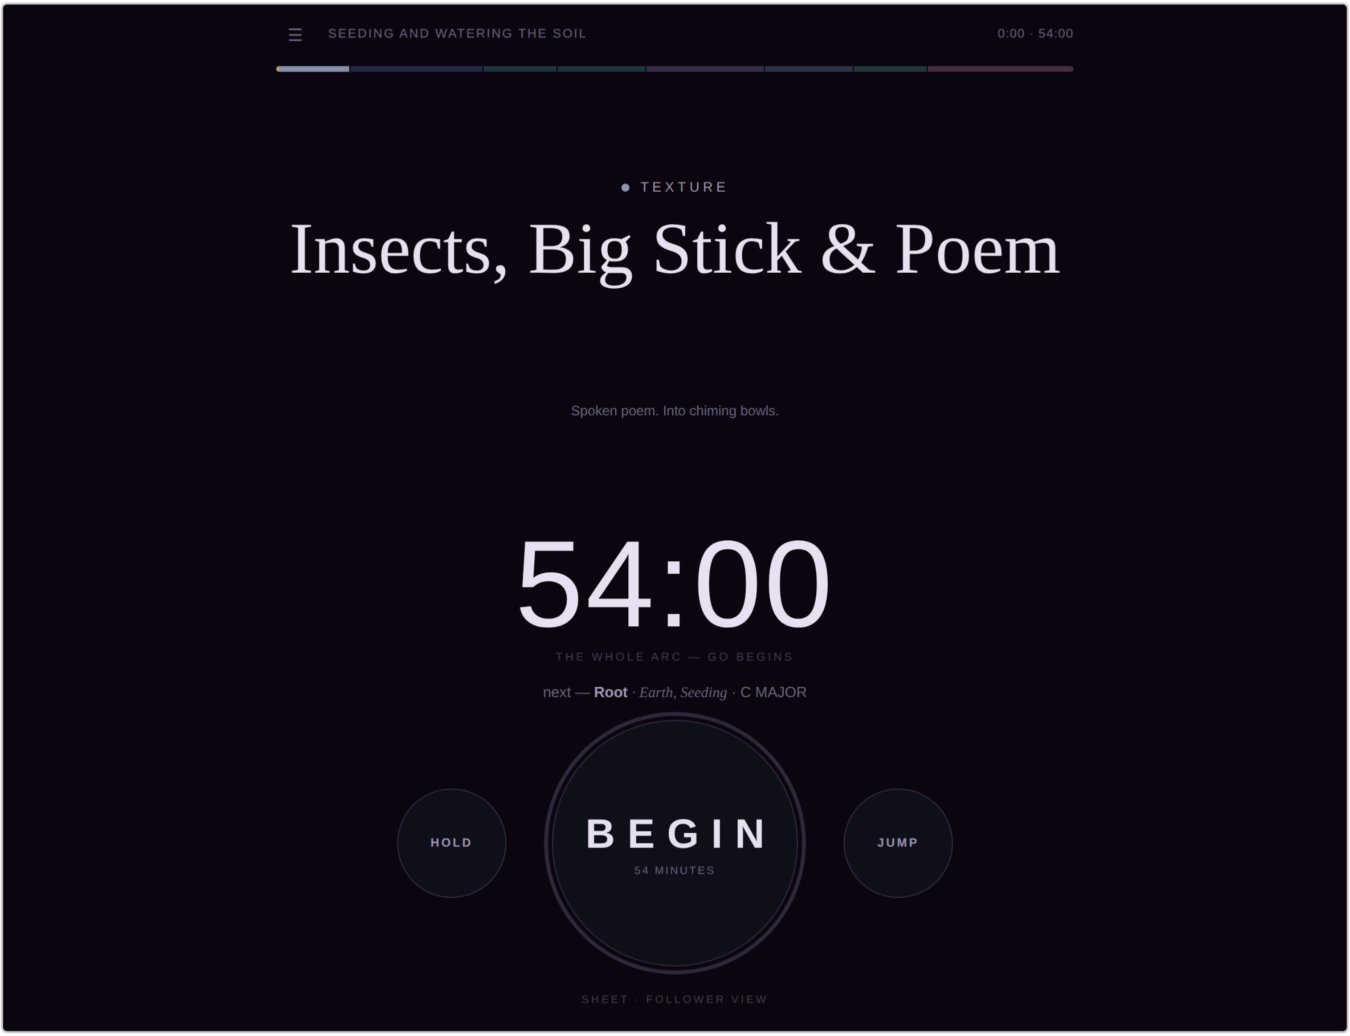

Perform. The cue clock runs the plan itself, so display-only is genuinely display-only — a leader with two mallets never has to touch it. GO advances early, HOLD freezes the clock, JUMP goes anywhere. The bowl sheet is right there, tiers carried by size. The whole interface drops to ember the moment the room goes live.

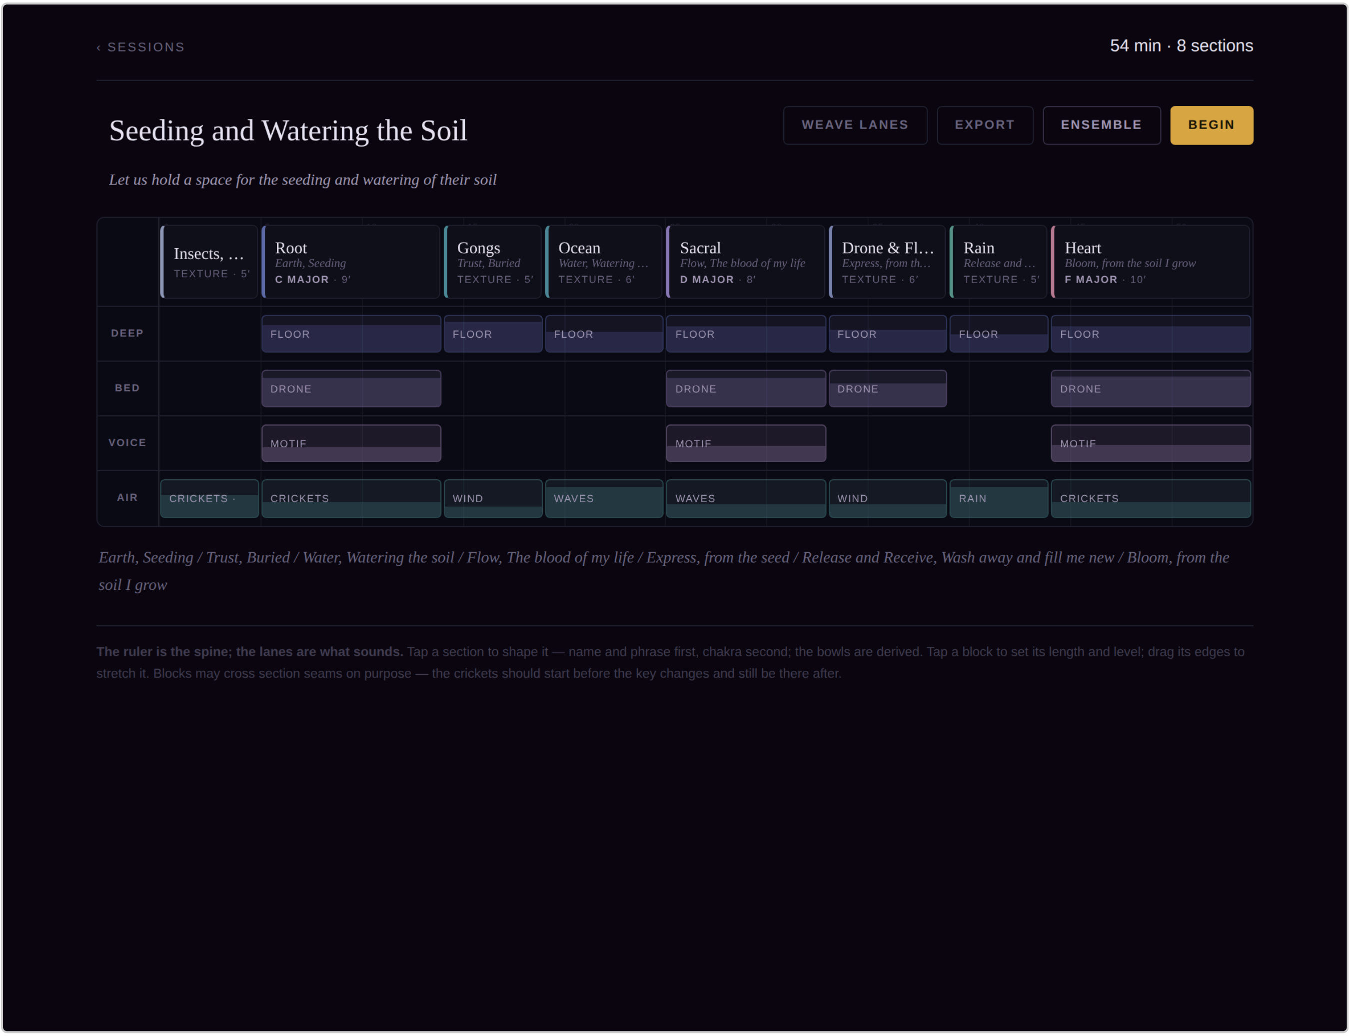

Shape. The marker ruler is the spine; four named lanes are what sounds. Blocks cross section seams by construction — the crickets start before the key changes and are still there after. The poem runs live beneath the timeline.

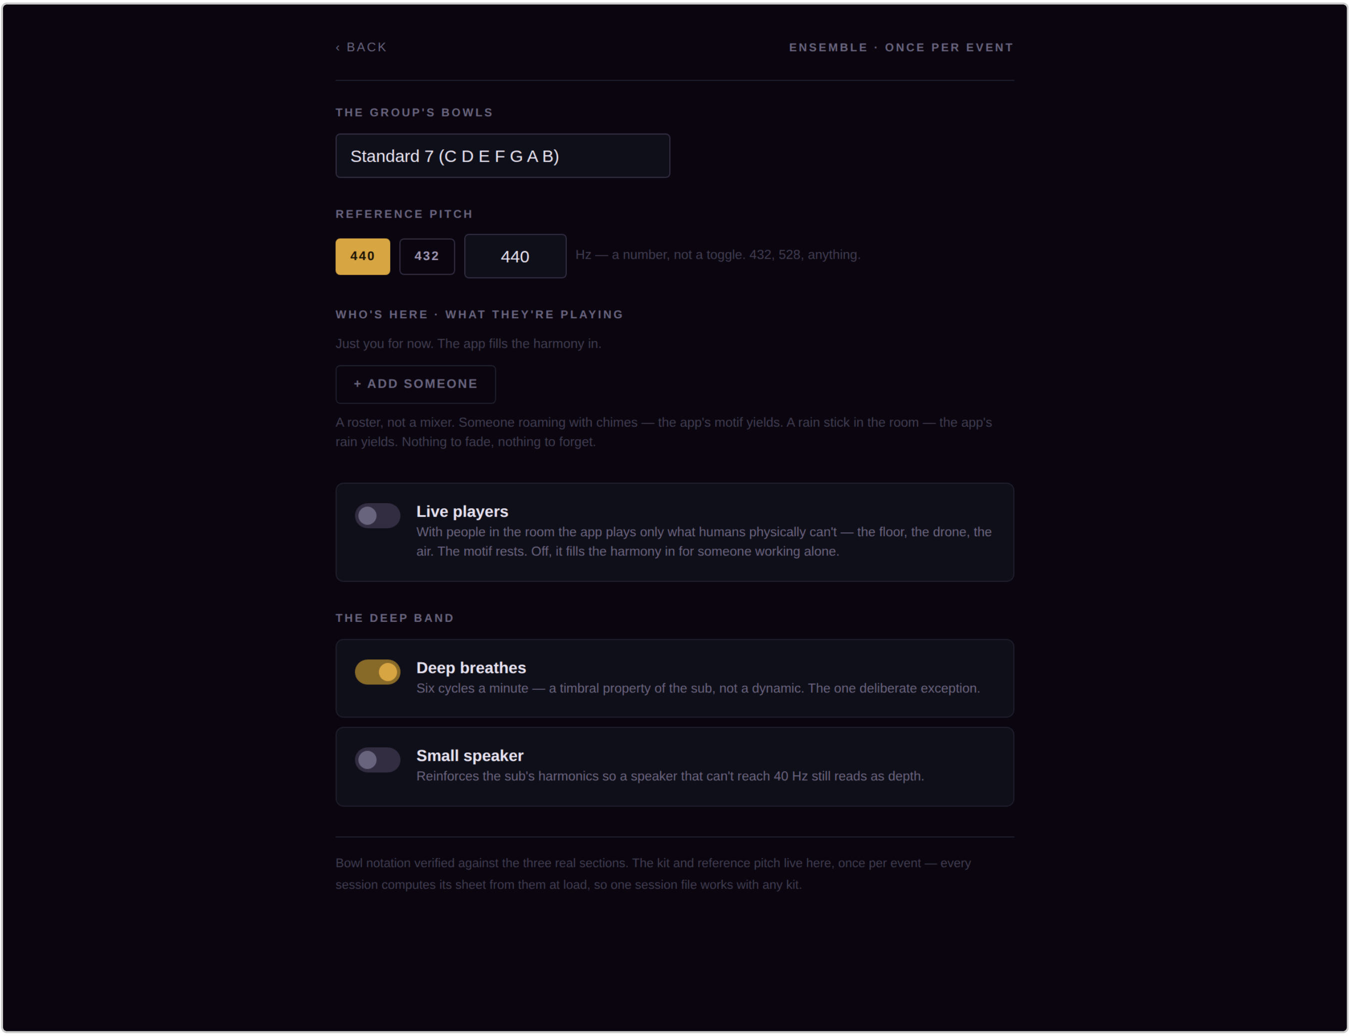

Ensemble. Once per event: the kit, the reference pitch as an editable Hz field, and a roster rather than a mixer. Note the copy — it teaches the reasoning without lecturing.

> **THREE CORRECTIONS CAME OUT OF BUILDING IT**
>
> Each surfaced only by implementing against the real fixtures, and each is now folded back into the brief. The demotion rule's scope — it applied to the leader's live chords, which provably brackets a bowl the real sheet puts in the free tier. `fadeIn` meaning two different things depending on how the playhead arrived. And a latent bug where Air's stored `fadeOut` was never populated, so every authored air exit silently fell back to 40 s.

Still unbuilt: multi-device sync, lights, live instrument inputs, the marketplace.

How the sound works

> ## The bowls own the mids and the top.
> ## The app owns the floor.
>
> Crystal bowls live roughly 200 Hz – 4 kHz and produce essentially nothing below that. That's physics, not technique. The app isn't backing music — it's the missing half of the spectrum.

There is **no PA and no microphones** at these events — bowls are acoustically loud on their own. So the app is the only thing needing amplification: one source, one speaker, no mixer, no feedback. The honest pitch is *your bowls and one speaker.*

THE FOUR BANDS

Allocate by **frequency band, not track count.** Amateur mixes turn to mud through frequency collision, not track quantity — bands can't collide by construction.

BAND	RANGE	CONTENT AND RULES
Deep	20–80 Hz	The sub. Nothing else is allowed down here.
Bed	low–mid	Drone always, chords optional. One drone only — two is mud by definition. Never a third.
Voice	mid	Rhodes motif, harp, chime. One tone every 9–26 s. Never a melody.
Air	1 kHz+	Crickets, waves, wind, rain. Two or three may coexist — they do in nature.

THE CARVING RULE

App content in the bowl range stays **sustained and static**. The bowls provide all the motion. If both move they mask each other, and the reflex fix — turning the pads down — makes them useless. Sparse and still in the middle, rich below and above.

THE FIVE LAWS

Structural, not conventional. These answer most future feature requests without a discussion.

1 **No pulse. Ever.**

Nothing is scheduled to a tempo. Sections are measured in minutes. No BPM, no metronome, no grid, no quantised cue launch. Conveniently this deletes the hardest problem in cue systems — there's nothing to quantise to.

---

2 **No third in the drone. Ever.**

Root, fifth, octave, ninth only. A third-less drone is a *permission structure* — it lets live players choose major or minor. Stating a third decides for the whole room. This is why tanpuras and shruti boxes are root-and-fifth and always have been.

---

3 **No automated dynamics.**

Nothing swells on its own. The arc is built from *arrangement* — a layer entering is a discrete event humans can play against; a slow crescendo is something they must fight or follow blindly.

---

4 **No competing melody.**

A motif layer and a live improviser in the same register is two soloists.

---

5 **Never bright by default.**

Brightness is the single variable deciding whether a pad is magic or unbearable, and it's asymmetric — slightly dark is pleasant, slightly bright is fatiguing within minutes and ruins a room. A tone control exists to *add* air for a full absorbent room, never to rescue a bad default.

> The best compliment the bed can get from an improviser is that they forgot it was there.
>
> Every instinct in music software pushes toward impressive. That instinct is the enemy of this layer.

KEY MIGRATION — THE SIGNATURE GESTURE

Someone holding a session in D taps F♯ and over about thirty seconds the entire room becomes F♯ — bed, texture, motifs, sub, all of it, seamlessly. Impossible acoustically, impossible with recorded material. It only works because everything is generated. **This is the moat in one move.**

- **True common tones don't move at all.** Matched by pitch class across the whole voice set, not positionally.
- The root octave is chosen nearest the previous root, so every voice travels the same small interval — C to B is a semitone down, not eleven up.

- Voices keep their **role** permanently, and levels travel with the pitch. Otherwise a voice landing on top arrives at root volume and screams.
- **Never glide everything simultaneously** — that's a portamento swoop, cool twice and nauseating on a floor of closed eyes.
- The sub moves last and slowest, and *crossfades* rather than sweeps when the move is large.

MEASURED BEHAVIOUR

PROPERTY	RESULT
Sub purity	2nd harmonic **−92 dB** relative to fundamental — a genuinely clean sine
Sub range	38–74 Hz. Below ~35 Hz the ear stops hearing pitch and starts hearing cycles as flutter
Sub modulation	**7.6%** — the 6/min breath and nothing else
Migration bound	No voice ever moves more than **5 semitones**; levels identical after every migration
Tone control	Monotonic 436 / 548 / 866 Hz across dark / flat / bright, and **0.10 dB** effect on the sub — dark must not mean more sub
Level-dependent tone	**0%** spectral drift across the full fader range

THREE BUGS WORTH REMEMBERING

Each was invisible without measurement, and each would have quietly ruined it.

The siren. Migration matched voices to pitches by nearest distance, ignoring role — so the bass voice got assigned the top pitch and swept 28 semitones up at root volume. A two-semitone key change produced a two-octave leap.

The limiter eating the sub. A 3 ms attack against a 40 Hz waveform (25 ms per cycle) tracks *inside* the cycle and turns a sine to grit — while pumping the whole mix at the sub's rate.

Level-dependent brightness. A post-fader send into a high-passed reverb means turning a band up adds proportionally more top, so the timbre changed as you moved the fader.

THE PALETTE

Yes	Drones, pads, sub, felt/Rhodes motif (Steven Halpern's *Spectrum Suite* is the reference timbre), harp and toy plucks, chimes, crickets, waves, wind, rain, rain stick, wave drum.
No — flute	Needs breath, tonguing and articulation. Synthetic flute is one of the most instantly fake sounds there is, and one bad voice cheapens the whole palette. Real sessions use a human flute player; leave it there.
Probably no — birds	Fast pitch-swept and formant-heavy, goes uncanny quickly.

SOUND WORK STILL OWED

1. **Foley must become world class.** The split is *procedural weather, sampled biology.* Rain, wind and waves are filtered-noise processes where synthesis genuinely *wins* — a sampled loop reveals itself over forty minutes, synthesis never repeats. Crickets, birds and frogs are behavioural and spectrally messy; synthesis loses. Sampling those breaks nothing, because **Air is unpitched and migration only binds pitched material.**
2. **Wavetables for the Bed.** `createPeriodicWave()` takes harmonic amplitudes and plays them at any pitch — sampled character, synthesised pitch, migration intact. Design a pad in a DAW, render a sustained note, extract its spectrum. **A wavetable is a few hundred bytes**, so custom timbres stay inside the preset-not-audio architecture: a curator's signature *sound* ships as JSON.
3. **Sub translation across speakers.** The sub sits from 41 Hz (E) to 73 Hz (D) depending on key. On a Bluetooth speaker rolling off at 80 Hz, 73 Hz is faintly there and 41 Hz is *absent* — so a migrating session would have the floor appear and vanish. Harmonic reinforcement must scale with how low the fundamental sits. **Never ask whether the user has a sub** — most don't know. Always reinforce at a level masked when the real fundamental is present.
4. **Evolving layers, scoped correctly.** With live players the bed stays static — they provide the variance. Solo, the app provides it. Implementation must be *layer entries timed to the section*, not free-running swells.

The domain facts

These came from real practitioners and real handouts, not from reasoning. They cannot be re-derived. This is the hardest-won material in the whole project.

CHAKRAS DETERMINE KEYS — NOT THE OTHER WAY AROUND

A standard seven-bowl set means **bowl number = chakra number = scale degree in C**, all three at once. That's why C, D and F dominate real sessions: Root, Sacral, Heart.

BOWL	NOTE	CHAKRA	INTENTION
1	C	Root	Security + Groundedness
2	D	Sacral	Creativity + Pleasure
3	E	Solar Plexus	Courage + Discipline
4	F	Heart	Love + Compassion
5	G	Throat	Connection + Expression
6	A	Third Eye	Intuition + Wisdom
7	B	Crown	Faith + Excitement

The app never asks for a key. It asks for a chakra and derives the key, with mode as an optional refinement — real sheets show both `C MAJOR` and `C MIXO`.

ON THE SCIENCE, HONESTLY

The chakra-to-Western-note mapping is a 20th-century Western construction — classical tantric sources don't assign pitches, and the competing systems (this one, the Solfeggio set, Cousto's planetary frequencies) don't agree with each other, which is the tell. 440 Hz is itself a 1939 committee decision, so 432 isn't wrong, it's differently arbitrary.

What *is* real is the effect, and the mechanism isn't mysterious: sustained low-stimulation sound, lying still, breathing slows.

So: **treat it as vocabulary, not mechanism.** The value of "Heart = F" isn't that F resonates with a heart — it's that everyone in the room already agrees, so ten people coordinate without a conversation. That's a shared protocol, and shared protocols are valuable regardless of physics. **The app labels; it never explains.** Show `HEART · F` and never a line of copy asserting what a frequency does to a body — which also keeps you clear of health claims.

THE BOWL SHEET NOTATION

Practitioners already hand out a printed sheet. It encodes four tiers in **punctuation, not words**:

Bold	the root of the section
Plain	freely available
(Parentheses)	sparingly
[Square brackets]	restricted — *the notes that contradict the stated mode*

Verified against three real sections. Every bracketed bowl is exactly the one that fights the key — computable, not a taste call:

SECTION	MODE	NOTATION	WHY THE BRACKET
ROOT	C mixolydian	`1 3 5 (2 4 6) [7]`	mixolydian wants B♭; bowl 7 is B♮
SACRAL	D major	`2 6 (3 5 7) [1, 4]`	D major wants C♯ and F♯
HEART	F major	`4 6 1 8 (2 3 5) [7]`	F major wants B♭

> TWO CORRECTIONS FOUND WHILE IMPLEMENTING
>
> **Bowls may state the third.** The no-third law binds the synth drone only — the humans are free, and the sheets put the third in the plain tier.
>
> **The octave bowl is only named when the doubling adds a non-root chord tone.** That's why C reads `1 3 5` and F reads `4 6 1 8`.

THE BOWL SET CONSTRAINS THE INSTRUCTION, NOT THE ENGINE

Pads and chords may use any pitch. The only requirement is that each section can name two or three bowls consonant against everything in it.

Rule: never leave the bowls with nothing to do. And weight *sustain*, not just interval — a bowl rings for twenty seconds, so a tension that would pass on a piano becomes permanent.

SESSIONS ALTERNATE PITCHED AND TEXTURAL

From two real handouts:

```
[texture] → ROOT (C) → [texture] → SACRAL (D) → [textures] → SPACE → HEART (F) → [texture]
```

Pitched sections carry chakra, mode and bowls. Textural ones carry neither — Ocean, Rain, Gong, Chimes, Insects. **Space is a programmed element in its own right**, not a gap. Both observed sessions ascend C →

D → F and close on a texture.

EVERY SECTION CARRIES A PHRASE

Read down a real sheet:

> # The rock / What grows / Will live / And shed / What's shed / Will burn / What's left / Is Home
>
> Under a title: LET IT BURN WHAT IT MUST. The session is one sentence broken across its sections — which is how the leader and the whole ensemble know what a section is *for*.

So the design flow asks **“what is this section about?”** before it asks anything musical. Name and phrase first, chakra second, bowls derived.

> **THE GAP, VISIBLE IN BALLPOINT**
>
> On the real handouts, the printed content is chakra, key and bowl numbers — the follower layer. The *handwritten* annotations are live-instrument chords and instrumentation — the leader layer. He is manually maintaining the leader/follower split in pen, because no tool separates them. That's the product, made visible.

THE TAXONOMY — THE INTELLIGENCE OF THE WHOLE PRODUCT

Packs, templates, the wizard and the design UI all derive from a canonical set of sound bath *functions*. This is the most important open item, and it has to come from studying real sessions.

Hypothesis to test: Arrive → Ground → Journey → Release → Return. Both known sheets fit it. Test against the next thirty.

Where the data is, easiest first: certification curricula (every training program had to name the functions in order to teach them); full session recordings on Insight Timer and YouTube, timestamped yourself; event and ticket copy; then real session sheets, which are richest and hardest.

Sampling beats volume. A thousand sessions from one scene tells you about that scene. A hundred across crystal bowl players, Himalayan bowl players, gong baths, yoga nidra, kirtan and breathwork tells you what's *invariant*.

How the app works

TWO MODES, DELIBERATELY UNEQUAL

	DESIGN MODE	PERFORM MODE
Used	a handful of times	every session
Posture	seated, planning	standing, hands full, dark room
Priority	second	**first**

Practitioners build somewhere between four sessions a year and one every few months, and repeat them freely — audiences never notice, because live cueing and live improvisers make every run different. **Variation happens at performance time, not design time.**

PERFORM MODE

Leader view — four things and nothing else: current section with its phrase, time remaining, what's next, and one enormous **GO**. Plus **HOLD** (stay here — the room isn't ready, a player is mid-phrase) and **JUMP** (go anywhere out of order). Without HOLD the app is driving instead of following, which is the failure that makes a practitioner resent it.

Follower view — this is already designed. It's the printed handout, one section at a time, on a phone.

PHYSICAL CONSTRAINTS, FROM REAL EVENT PHOTOGRAPHS

Nowhere to put a phone inside the bowl arc — bowls are packed edge to edge with candles in the gaps. No free hand: a mallet in each. Screens are read at about four feet, at an oblique angle, in low light. **Rooms are washed warm amber**, so warm section colours vanish and cool ones survive. Screens must drop to dim amber and deep red once a session starts — a bright white phone in that room is disrespectful and everyone knows it.

And **fade times are ergonomic, not aesthetic**: the crossfade is the window in which a human sets down one bowl and picks up another. Thirty seconds is indulgent as a musical choice and exactly right as a physical one.

DESIGN MODE — A MARKER RULER OVER FOUR FIXED LANES

The layout, precisely:

- **Across the top, a marker ruler.** Section markers placed by time, the way arrangement markers work in a DAW. Each is chosen from the **canonical function language** — Welcome, Ground, Journey, Release, Return — and freely relabelled for a particular session (*"The rock"*, *"What grows"*). The function sets the chakra, mode and band defaults; the label is the poem.
- **Below it, four named lanes** — Deep, Bed, Voice, Air. Named after what they *are*, not after track numbers, so nobody has to remember what track 3 was.
- **Blocks stretch** to the length you need, and live on the lanes rather than inside the sections.

That separation is the important part. **Layers must be able to span section boundaries.** If every band changes at the seam, every transition is a hard cut — the crickets should start before the key changes and still be there after. So the markers are the *spine* and the lanes carry what's actually sounding, with their own starts and ends.

A worked example of what someone is actually building:

```
15 min — Ground (C) · drones + waves  →  8 min — Journey (D) · drones + crickets  →  12 min — Release (F) · drones + rain
```

…with the crickets deliberately overlapping the boundary into the third section, so the room changes without anything announcing it.

	GROUND **The rock** 15 min · C	JOURNEY **What grows** 8 min · D	RELEASE **What's left** 12 min · F
DEEP	sub — whole session		
BED	C drone	D drone	F drone
VOICE		Rhodes motif	
AIR	waves	crickets ⟶	rain

The crickets deliberately cross the D→F boundary. Sections change; the texture carries over. That overlap is the difference between a room that breathes and one that steps.

WHAT KEEPS IT FROM BECOMING A DAW

The lanes are fixed at four and never grow. That's the entire difference between this and Ableton's arrangement view — unbounded tracks is where complexity explodes. Four lanes, fixed roles, one thing at a time, blocks that stretch. Closer to a lighting console than a DAW, and that limitation is the thing that makes it usable by someone who cannot compose.

Tap to set length as the primary interaction, drag as secondary — dragging block edges precisely with a finger is the thing that would make it feel fussy. Snap to whole minutes.

ENSEMBLE SETUP, AND ONE SWITCH

A **roster, not a mixer**: who's here and what are they playing. Someone roaming with chimes → the chime layer mutes automatically. One setup screen adapts a single design to three people or ten without anyone touching a fader.

Live players mode — flip on and the chord layer mutes, the motif layer pulls back, and the app plays only what humans physically can't. Flip off and it fills in the harmony for someone working alone. That single toggle encodes the entire philosophy: with people in the room, get out of the way; without them, be the ensemble.

MULTI-PRACTITIONER SYNC

Groups of four to ten are real — synchronised off one person who knows the structure, with hand signals for sections. Which reframes the product: **it's a conducting tool that happens to make sound.**

- **One device makes all the sound; every other is a display.** Syncing audio across devices is brutal — clock drift, jitter. Syncing display state is trivial, and a quarter-second of lag on "we're in section 3" is invisible.
- Broadcast the **active pitch set**, not per-person assignments — each phone filters against the kit its owner declared once. No roster to maintain, self-corrects when someone doesn't show.
- Ship suggested **hand signals** too. Fewer screens in a dark room — and if the app defines the signals, the app becomes the standard. That's a social moat, not a technical one.

ARCHITECTURE DECISIONS TO MAKE NOW

Event broadcaster	The cue engine emits `section:change` , `key:migrate` , `fade:start` from day one. Follower phones, lights and projection are all just subscribers. Nearly free now, painful to retrofit.
Reference pitch	An editable Hz field, never a 432/440 toggle. Then 432, 528, the whole Solfeggio set and every future request fall out of one control.
Packs are presets	Rendered audio can't migrate key with the room, so the first purchase would break the signature gesture. Sessions are parameters — kilobytes, not megabytes. **Templates** ***are*** **packs**, same file type.
Lights never touch audio	Analysing a signal is fuzzy and dumb when the app already knows it's entering section 4 in F.

Marketing

> ## The hero is always the person. Never the software.
>
> The bowl player who becomes someone who can lead ten practitioners. The meditation teacher who becomes someone who designs a whole experience.

THE LINE

> ## One violin player becomes a symphony.

Concrete, about them rather than about you, and the whole promise in six words. Keep it.

THE FILM — A SHOT LIST

Two minutes about a Thursday that turned into something, with **three cuts of a screen in it.** That ratio is also the product philosophy: the technical layer disappears.

01 Thursday. A text lands — musicians you love are passing through.	02 A beautiful room with plants. Tea being poured. Tarot cards on a low table.	03 Two friends leaning over a laptop. Fifteen minutes of structure. *Screen cut one.*
04 Phones syncing — the pulse. *Screen cut two.*	05 Bowls coming out of a canvas bag. Candles being lit. Mats laid down.	06 People arriving who didn't know this existed on Monday.
07 Eighty seconds of the room. Live players. Light. Sixty people on the floor.	08 **The insider shot:** a hand raised mid-session, and the room changes. No caption. *Screen cut three.*	09 After. Someone stacking mats, smiling. **“We made it happen.”**

Shot 08 is the one that sells it to a practitioner rather than a general audience. Anyone outside the world reads it as atmosphere; anyone inside it sits up, because they know exactly how hard that is to do with ten people and no click track — and that they can't currently do it.

Keep the pronoun in the closing line. *We* is the people in the room, not the app.

WHAT THE TEA AND TAROT DETAIL ACTUALLY TELLS YOU

It's product input, not set dressing. The evening starts long before the sound does — arrival, settling, talking — and that stretch still needs to sound like something. Nobody would think to compose it, which is exactly why a template should handle it by default. Most amateur sessions probably fumble the first fifteen minutes because nobody thinks of it as part of the piece.

RULES FOR ALL COPY

Never say	“Your audience can't tell the difference.” True, and it reads as contempt for the craft to the exact people whose endorsement you need.
Never say	“AI.” In this market it signals inauthentic, extractive, machine-made spirituality. Call it planning the arc, or just ask the questions.
Never say	Anything about what a frequency does to a body. Labels only. Health claims are a different legal posture.
Do say	The sound design was never the art — leading the room is. This removes what was blocking you from what you're good at.

AESTHETIC DIRECTION

Modern and hip, but wellness. Draw from yoga, tea ceremony, ceremony generally — not from music software and not from spa clichés.

Steal	Lumen by Kompose Audio — the four-strip anatomy, near-monochrome restraint, a thin high-contrast serif over a plain functional sans. That pairing says *beautiful object* and *actual tool* at once.
Invert	Its value. Lumen is a light UI for a bright studio. This runs in a candlelit room, so light UI is disqualifying.
Kill	Fantasy sample-library naming ("Aeonis Pad"). Tab bars. Keyboard graphics. Anything reading as music production software.
Palette	Near-black with a violet cast; atmosphere in muted violet and silver; **warm amber reserved exclusively for "this is live."** Section coding runs cool — indigo, violet, teal, soft rose — because the rooms are already warm.
Never	Near-black plus an acid accent ("AI dark mode"). Chakra rainbow at full saturation on a screen — half those hues vanish in candlelight.
Type	Thin high-contrast serif for identity only. A clear sans for everything operational, at sizes that feel slightly too large on a desk — because they'll be right on a floor, at four feet, in the dark.

NAMING

Candidates: **Bed**, Bath, bathsound.io. *Atmos is unusable — Dolby owns it in audio.*

The constraint here is the **opposite** of Pedal's. Pedal's main risk is discovery, so it needs a searchable name. This sells hand-to-hand inside a community, so memorability and meaning matter far more than search volume — which makes a generic English word much cheaper here.

Bed carries the most meaning: the musical bed, the thing people are lying on, and the sense of being held, all at once. "Bath" collides with bathtubs everywhere and adds nothing. "bathsound.io" inverts a compound everyone knows in the other order, so people will type soundbath.io and land somewhere else.

Business

Price like gear	This community buys $800 bowls without blinking. They're not price-sensitive; there's just nothing to buy. Consumer-app pricing is badly wrong here.
A rig, not an app	Nobody pays four figures for software they open on a phone. They do pay it for a *system*. The sub is what makes the price legible.
Rent the sub	The strongest hardware move available. A practitioner needs a sub four or five times a year — the textbook definition of a rental good. It solves the "half your users can't hear the differentiator" problem directly, puts the app in a room with an audience every time, and $100 against an event grossing thousands isn't a decision. No manufacturing, no returns, no freight. It's also the cleanest market research: whoever rents twice is a real customer.
Hardware sequencing	Recommend → affiliate → rent → bundle. **Never manufacture a speaker** — heavy freight plus damage economics eats small companies. If anything ever carries the brand physically, make it the LED bowl pads.
Group product	One leader licence, N follower devices. Nine practitioners experience the app from the inside every session, and some fraction will want to lead their own.
Packs, not subscriptions	Four sessions a year means packs are closer to buying another bowl than to a subscription — occasional, considered, a genuine expansion. Real revenue, wrong model as churn.
Don't raise	A few thousand practitioners at a few hundred dollars is a genuinely good business and a genuinely terrible venture business. That's *why* the gap is empty — nobody funded is coming for it.

READ THE FIRST SALES CORRECTLY

Ten warm sales in Ohio validate your relationships, not a market. The real signal is buyer eleven — someone who found it cold and paid anyway. Ten warm sales are still absolutely worth doing.

THE MARKETPLACE, EVENTUALLY

Curator identities with named packs, each containing coherent *sections* rather than one track:

```
ATMOSFLY → Venus → Opening (C) · Grounding · Unwinding · Release · Restore · Return
```

- **Packs are presets.** Kilobytes, no CDN, 432 and 440 are a field not a product variant, and purchased material migrates with the room. Splice ships gigabytes; this ships JSON.
- **Wavetables mean packs can have a genuine sonic identity** — not just distinctive settings — and still be kilobytes.
- Cold-start fix: the house makes the first inventory. Three or four in-house curator identities is a populated shelf on day one without pretending.
- Then **local legends** — the biggest practitioner in each regional scene writes a pack and becomes an affiliate. That's how it starts, not with a marquee signing.

What we are not building

A DAW	Infinite choice is the wall this exists to remove.
A listening app	The category is packed. This is for the person *leading*.
Anything called AI	The wizard is a decision tree — instant, offline, free, can never say anything strange.
Health claims	Labels only.
Audio in packs	It would break key migration on first purchase.
Manufactured speakers	Freight and damage economics.
Audio sync across devices	Unnecessary — one device makes sound, the rest are displays.
Venue lighting rigs	LED bowl pads are already in the room and are the right v1.
A pulse, anywhere	Law 1. Ecstatic dance wants one; that's a different product.
Birds and flute	Both go uncanny synthesised, and one fake voice cheapens everything.

Roadmap

Phase 0 ✅	**Validate the sound.** Done. Engine built, measured, assessed as working. Sub confirmed as a genuine differentiator.
Phase 1 ✅	**The sheet generator.** Built. Not yet handed to real bowl players — that test is outstanding and costs nothing.
Phase 2	**Templates + perform mode.** Cue list with HOLD/NEXT/JUMP, follower view. Performable without a designer existing.
Phase 3	**Design mode.** Four-lane stretchable blocks, the wizard, ensemble config.
Phase 4	Pack format hardening, in-house curator packs.
Phase 5	Lights (LED pads), live instrument inputs through the shared reverb, marketplace, rental programme.

LIVE INPUTS ARE WORTH PULLING FORWARD

Not future, and not hardware — an iPad plus a ~$100 USB interface is already two inputs. The win is **shared reverb**: run the live violin through the same reverb as the generated bed and the player stops sounding like they're on top of the track and starts sounding like they're inside the piece. Right now those two things live in different spaces and every ear in the room knows it without being able to name it.

PART TWO

Pedal

A drone-first practice tool for improvising musicians. Shipped and live — one self-contained HTML file, no dependencies, works offline.

Why it exists

> Opening a DAW puts you in producer mode.
> This keeps you in improviser mode.
>
> Drones are how musicians practice, *and* drones are how you get into a flow state. The enemy of flow is the twenty minutes of setup and decision-making a DAW demands before a single note happens.

Speed of setup is the product, not a feature of it. Every design decision gets measured against one question: does this get someone playing sooner, or later?

And most practice happens **without bars.** An improviser working on ear training wants one chord droning forever while they feel out where the tones sit. Almost everyone building this makes a chord player and bolts a drone on as a degenerate case — one chord, infinite length. Inverting that is the whole idea, because tonic-relative hearing is the actual skill and a drone is the only way to isolate it.

WHY A PHONE GENUINELY WINS

DRONES HAVE NO LATENCY REQUIREMENT

Bluetooth's 150–200 ms lag makes a metronome unusable and a drone completely unaffected. So "pull it up, Bluetooth to whatever speaker is in the room, hit play" isn't a compromise — for this one use case it is strictly better than a hardware rig.

iReal Pro	Clumsy at holding one thing forever. Its moat is a library of thousands of user-entered charts, not its software.
Tanpura apps	Timbrally committed to Indian classical. Perfect inside that tradition, useless if you want a neutral pad.
Hardware workarounds	Hold a pad and hit sustain on a looper, or carry a freeze pedal. All worse than a phone in your pocket.

WHO IT'S FOR

A working improviser who has already made ghetto versions of this — exporting 30-second drones out of Logic and looping them on their phone for a decade. They don't need to be taught theory. Concretely: anyone working intonation against a reference (horn players, singers, strings, bass), the very large population of guitarists who want something to noodle over, and teachers who would send a link.

What's built

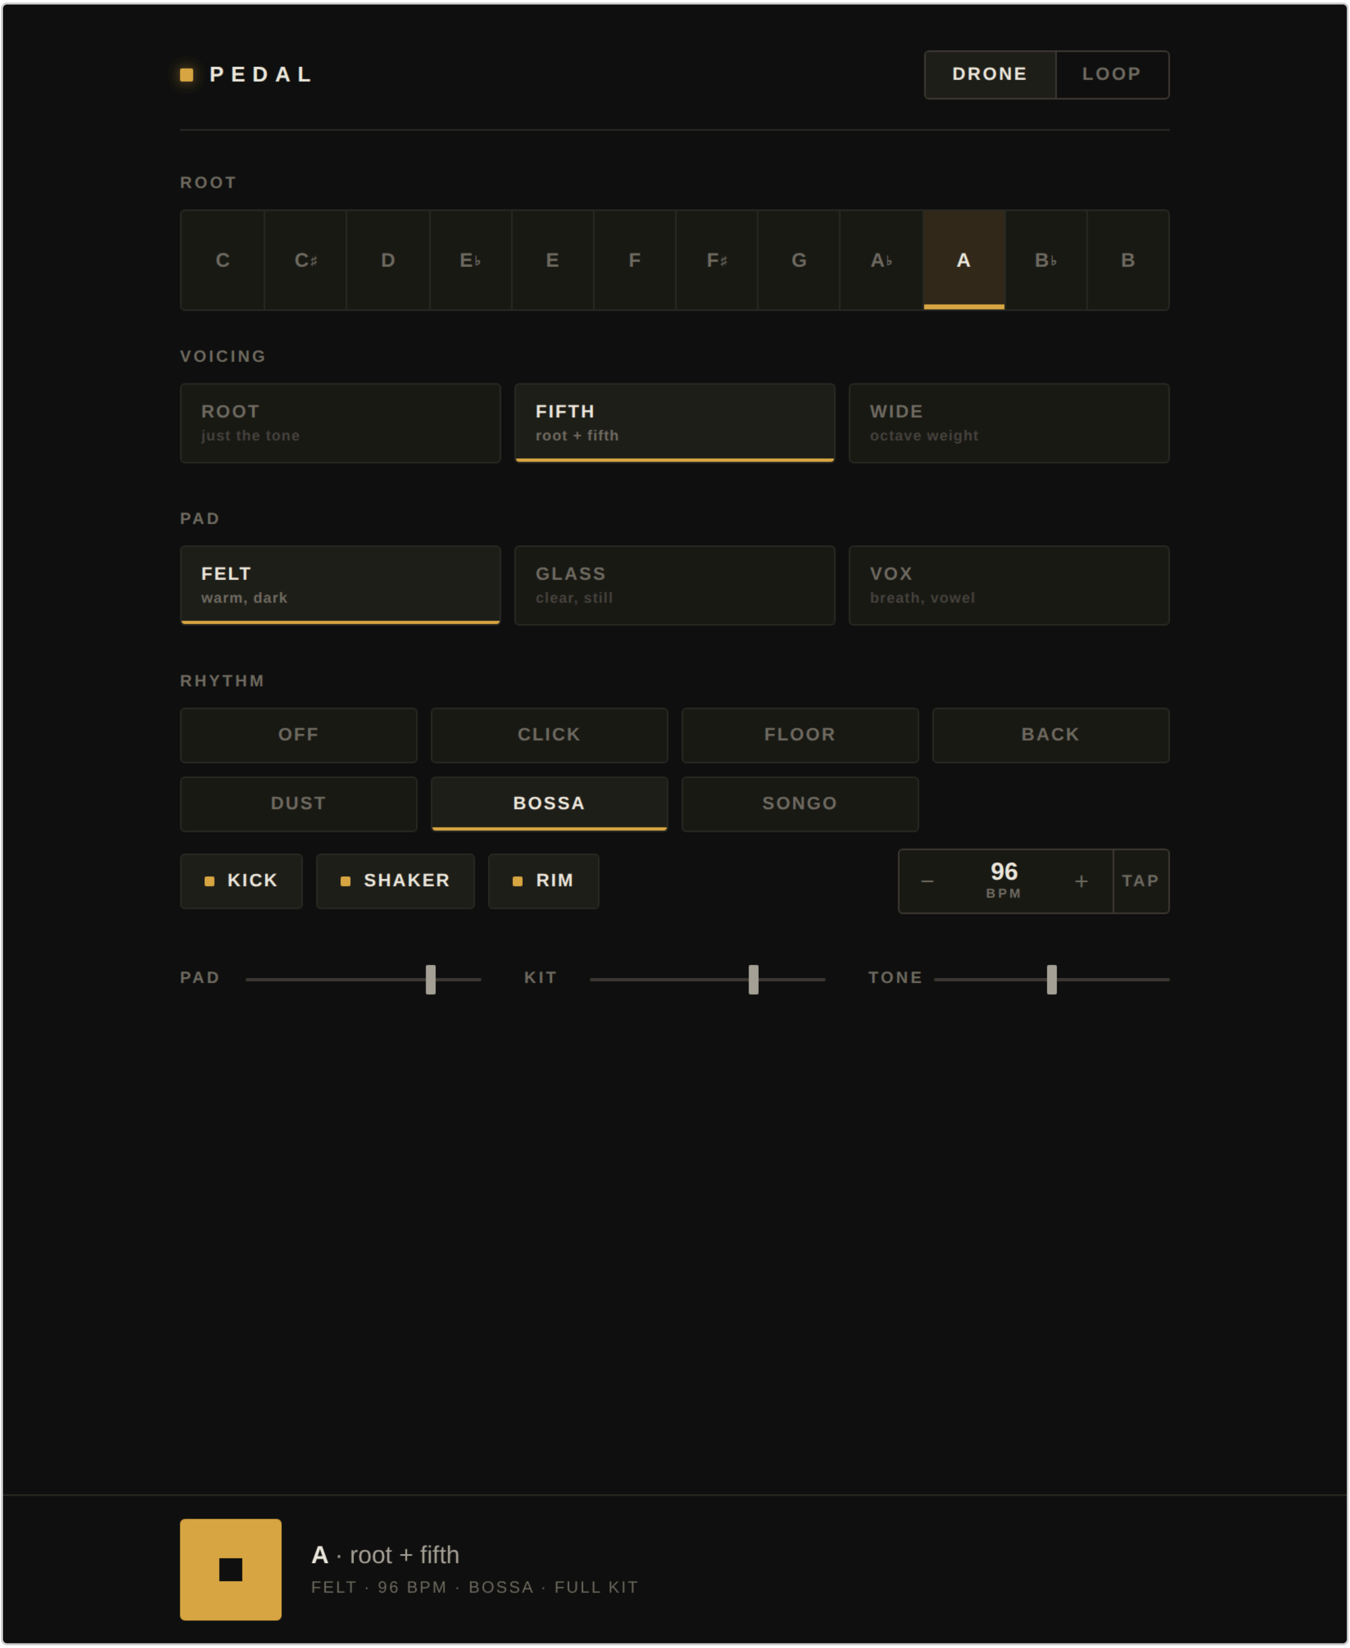

Drone mode — the default. Twelve roots, three voicings, and no bar counter or measure grid anywhere. Changing root while running glides rather than restarting.

Loop mode — blocks sized proportionally to duration, live playhead.

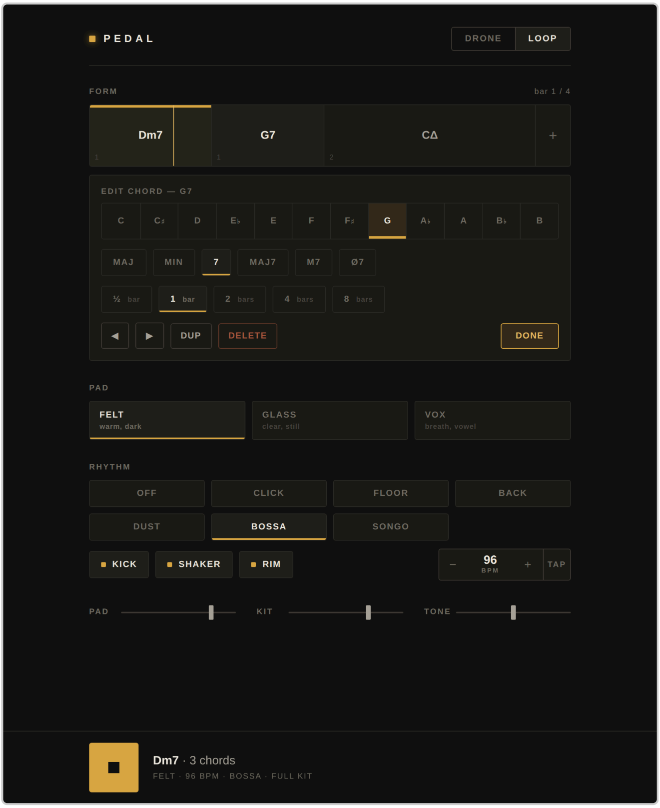

Tap a block to edit. Add stays open and pre-advances a fourth for fast entry.

SOUND

PAD	WHAT IT IS
Felt	Soft-detuned saws under a slowly breathing dark lowpass. Juno-blanket territory.
Glass	Low-index two-op FM plus sine partials at the octave and twelfth, each blooming on its own slow cycle. Nothing above a sine, so harshness is structurally impossible.
Vox	Detuned triangles through fixed vocal formants with a drifting second formant, so the vowel slowly turns.

Detune never exceeds 3.5 cents and no LFO runs faster than 0.1 Hz — so there is no wobble over an hour.

KIT	WHAT IT IS
Kick	Clean sine drop, 3 ms soft attack. Neutral and digital — not acoustic, not an 808, not a clicky EDM kick.
Shaker	Banded noise whose filter centre wanders per hit, so a long run never sounds stamped. Tempo-independent envelopes.
Rim	Three tuned resonators struck by a 4 ms noise chirp.

GROOVE	WHAT IT IS
Click	Straight per-beat metronome, accented downbeats
Floor	Four-on-the-floor, rim answering 2 and 4
Back	Backbeat, kick on 1 and 3 with an and-of-4 pickup
Dust	Lo-fi, swung sixteenths, humanised timing and velocity
Bossa	**True bossa clave** (1, 2&, 4 | 2, 3&) over the surdo foot pattern
Songo	**Son clave 2-3** (2, 3 | 1, 2&, 4) with tumbao kick anchors

Any kit voice mutes out of any groove — strip to just shaker, shaker and rim, whatever. That's the only customisation that matters.

CONTROLS

- **Tone** — a tilt EQ, two opposing shelves pivoting at 900 Hz, ±6 dB, so dark means *more body* rather than merely less treble. Verified monotonic: 181 / 242 / 366 Hz across dark / flat / bright.
- Pad and Kit faders. BPM by stepper, vertical drag, or tap tempo.
- Space toggles play; arrow keys walk the root chromatically.
- Screen wake lock. Full state persisted locally.

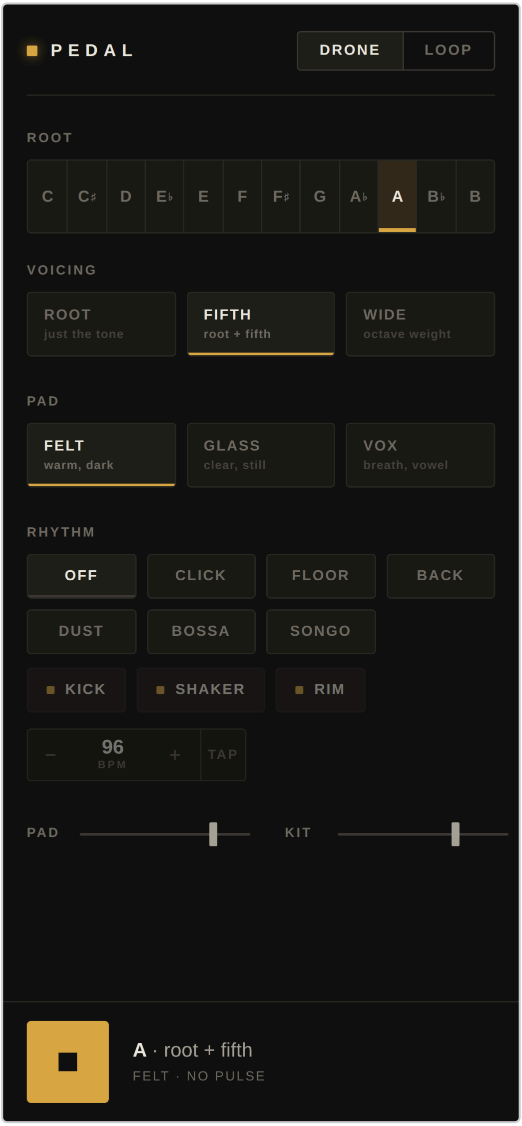

On a phone — two taps to sound.

What we are not building

A tune library	Unwinnable — iReal Pro's moat is thousands of user-entered charts. And it isn't the wedge; the drone side is.
Notation, diagrams, explainers	The user already knows. This is an instrument, not a teacher.
Accounts, sync, cloud	A URL plus localStorage covers the entire need.
More grooves	Six curated beats twenty. **A groove menu is a decision, and decisions are producer mode.** Adding one means cutting one.
A mixer, effect racks, automation	That's a DAW — the thing being escaped.

THE ONE OPEN QUESTION

Nobody has heard it on real speakers. It was built and metered, never listened to. The whole differentiator is "good enough that you'd leave it running for an hour," and that's unverified. The failure mode is *fatigue*, which doesn't appear in a thirty-second audition — it appears at minute twenty.

Roadmap

1. The listen	Priority zero. Every item below is worth less than a good-sounding pad. Known carry-overs from the Bed engine: level-dependent brightness (post-fader send into a high-passed reverb) and sub-range limiter distortion (3 ms attack against a 22 ms waveform).
2. URL state	`…/looper/#A/fifth/felt/bossa/96` and a student lands on exactly that setup. No account, no backend — **the URL is the save file.** This targets the real risk: the problem was never that nobody wants it, it's that nobody finds it, and "sounds better" survives neither a screenshot nor a feature list.
3. Capture	One button records everything — drone, kit and mic — while you play. Multiple takes, saved simply. **No latency requirement at all**, because nothing is layered back in real time. The value isn't fidelity, it's *"what was that thing I just played?"*
4. Wavetables	Sampled character, synthesised pitch, a few hundred bytes each.
5. Section markers	A 32-bar form needs visible A / B / bridge structure, not 32 identical blocks.
6. A darker fourth pad	Felt is the workhorse; Glass and Vox are both fairly bright.
7. Progression paste	`Dm7 G7 Cmaj7 | Am7 D7 Gmaj7` parsed. Two minutes becomes five seconds, without competing on catalogue.

WHY TRUE LOOPING IS HARD, AND CAPTURE ISN'T

Overdub looping collides with the one thing that makes this work on a phone. Drones tolerate Bluetooth latency; loops do not. Record over a playing drone and what you heard arrived 150–200 ms late, so the take lands that far behind — and every overdub compounds it. Fixable only with real calibration: play an impulse, capture it back, subtract the measured round trip.

Two more phone traps: browsers default `getUserMedia` to speech processing (echo cancellation actively ducks your instrument whenever the drone plays, noise suppression eats sustained tones, AGC pumps — all must be explicitly disabled), and internal mic plus speaker is a feedback loop with a sustaining drone as the worst possible case.

So looping ships as a mode you *enter*, with a stated setup. The front door stays "open it, press play."

NAMING

Working name **Pedal**. Candidates raised: DroneFlow, FlowSound.io, DroneLoop.io.

One caution: **"drone" means UAV to the general internet.** Search and ad keywords are dominated by quadcopters, which is a real discoverability tax on a product whose main risk is already discovery. A tradeoff, not a dealbreaker — but a decision, not an accident. Of the three, **DroneFlow** is strongest because "flow" encodes the value proposition rather than the mechanism.

PART THREE

The landscape

Researched, not assumed. Three crowded categories — and none of them is this.

CATEGORY	WHO'S THERE	WHAT THEY ACTUALLY DO
Consumer listening	The Sound Bath App, PAUSE (Sara Auster), Insight Timer, Sound Bath Journeys	Packed. All of it is *lie down and receive* — content delivered to a client.
Business software for healers	Flowdara, Breely, Acuity, ZenPass	Real category. Somebody already noticed these practitioners pay for tools — they built for the calendar, not the room.
Generative engines	**Wotja**, Endel	Wotja is the closest thing and ships with **130+ engine parameters.** Endel went the other way: fully automatic, biometric, zero authorship, made for earbuds not a room.

> Nothing exists for the fifty minutes a practitioner is standing in front of sixty people.
>
> Not failing. Empty.

Wotja's parameter count is the entire thesis in one object: the technology has been solved for years and packaged for exactly the person who isn't the customer.

Why it's empty: everyone chased listeners, rationally, because that market is millions and facilitators are thousands. A few thousand practitioners at a few hundred dollars is a genuinely good business and a genuinely terrible venture business. **Nobody funded is coming for this.** It requires fluency in both bowls and code, which is an unusually specific keyhole.

STRATEGIC CONSEQUENCE

Build lean, sell direct, don't raise, and never let anyone reframe it as a consumer play — that's the version that forces a pivot to listeners and becomes the eleventh app in category one.

FOUNDER FIT

The credential stack isn't the point. The point is that this product needs credibility in **two mutually suspicious communities** — wellness thinks tech is extractive and cheapens something sacred; music-tech doesn't take sound healing seriously at all. Almost nobody speaks both languages natively.

AND THE MATCHING RISK

A musician building for non-musicians consistently fails one way: **by adding flexibility.** It feels like generosity to the builder and reads as a wall to the user. The discipline is to be a *dictator* about timbre, tuning, fade curves and palette — they cannot evaluate those and shouldn't be asked to — and almost *cowardly* about the interface. Whatever seems like the minimum adjustable, it's less than that.

APPENDIX

Where everything lives

So this is never trapped in a chat again.

THE REPOSITORY

`adamelfersmusic-web/AdamVaultOS` — branch `claude/practice-looper-improvisers-naaagy`, merged to `main`.

PATH	WHAT IT IS
`sound-bath-app/VISION.md`	Bed, end to end — why, avatar, app, sound, business, branding, open tensions, roadmap
`sound-bath-app/BUILD-BRIEF-V1.md`	What a builder builds: v1 scope, domain facts, the bracket algorithm, fixtures, explicit non-goals
`sound-bath-app/IDEAS-LOG.md`	Every idea from the originating conversation, attributed, plus a table of corrections that moved the thinking
`sound-bath-app/FABLE-PROMPT.md`	The handoff prompt — deliberately short, points at the docs
`sound-bath-app/prototype/bed.html`	The four-band engine
`sound-bath-app/prototype/sheets.html`	The handout generator
`sound-bath-app/prototype/fixtures/`	Both real sessions as JSON
`public/looper/index.html`	Pedal — live at `/AdamVaultOS/looper/`
`public/looper/SPEC.md`	Pedal, end to end

`sound-bath-app/` sits outside `public/`, so it is intentionally not published. The two products are already cleanly separated — that folder lifts into its own repository with no untangling.

IF YOU'RE PICKING THIS UP COLD

1. Open `prototype/bed.html` on a real speaker. Ten minutes. Everything downstream is a judgement about that sound.
2. Read `VISION.md` §0.5 for what actually exists versus what's still imagined.

3. The **taxonomy** is the biggest unfinished thing and the one nobody can shortcut. Twenty to thirty more real session sheets and it hardens into the product's brain.
4. Hand `FABLE-PROMPT.md` to a builder for perform mode. Add two lines: build as a web app for now, working name *Bed*.

THE FOUR QUESTIONS THAT DECIDE THE SOUND

Timbre is tunable forever. Architecture isn't. So the feedback that matters is structural:

1. Is there room in the middle for bowls? *Answered: yes.*
2. Does migration land as a change of place or as an audible effect? *Answered: place.*
3. Is four bands the right split — anything missing or redundant?
4. Is anything wrong *in kind* rather than in degree? That's the only genuine binary; everything else is tuning.
